MONASH University

ETC5512: Wild Caught Data

Combining Australian census and election data

Lecturer: *Kate Saunders*

Department of Econometrics and Business Statistics

ETC5512.Clayton-x@monash.edu

Week 6

Today you will:

- look at the ABS geographical boundaries for the 2021 census
- integrate data from different sources (census and election) to make exploratory inferences

Coding Perspective:

- Further expand our understanding of how to read and use spatial data in R
- Better understand how spatial data is organised
- Learn how to intersect two spatial objections
- Practice re-projecting maps

Recall 2021 Federal Election Data

```
library(tidyverse)
library(sf)
aec_map <- read_sf(here::here("data/vic-july-2021-esri/E_VIC21_region.shp"))
votes <- read_csv("https://results.aec.gov.au/27966/Website/Downloads/HouseDopByDivisionDownload-27966.csv", skip = 1)

electoral_winners = votes |>
  mutate(DivisionNm = toupper(DivisionNm)) |>
  filter(Elected == "Y") |>
  select(PartyAb, DivisionID, DivisionNm, Elected) |>
  distinct()

aec_map = aec_map |>
  mutate(Elect_div = toupper(Elect_div))

winners_with_map = electoral_winners |>
  left_join(aec_map, by = c("DivisionNm" = "Elect_div"))

aus_colours <- c(
  "ALP" = "#DE3533", "LNP" = "#ADD8E6", "KAP" = "#8B0000", "GVIC" = "#10C25B", "XEN" = "#ff6300",
  "LP" = "#0047AB", "NP" = "#0a9cca", "IND" = "#000000"
)

ggplot(winners_with_map) +
  geom_sf(aes(fill = PartyAb, geometry = geometry)) +
  scale_fill_manual(values = aus_colours)

#UPDATED CODE - previous variable assignments from past weeks
winners_fix = winners_with_map
auscolors = aus_colours
```

Recall 2021 Federal Election Data

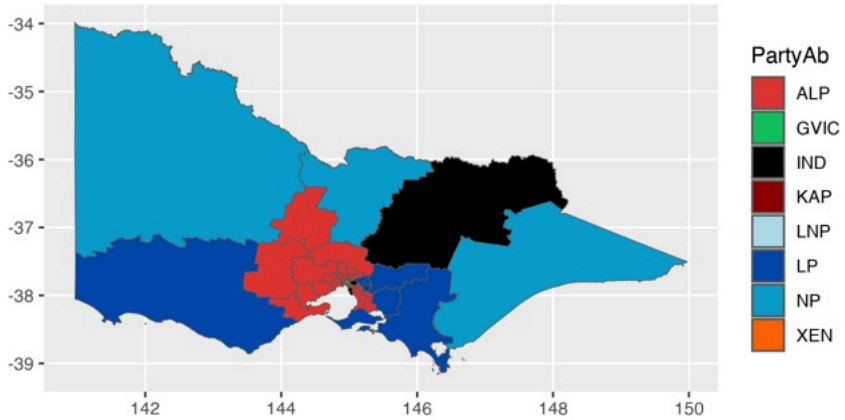

There are two sources of data:

1. Electoral boundary
2. The votes for candidates in each electorate

Recall 2021 ABS Census Data

- DataPacks https://datapacks.censusdata.abs.gov.au/datapacks/
- GeoPackages https://datapacks.censusdata.abs.gov.au/geopackages/

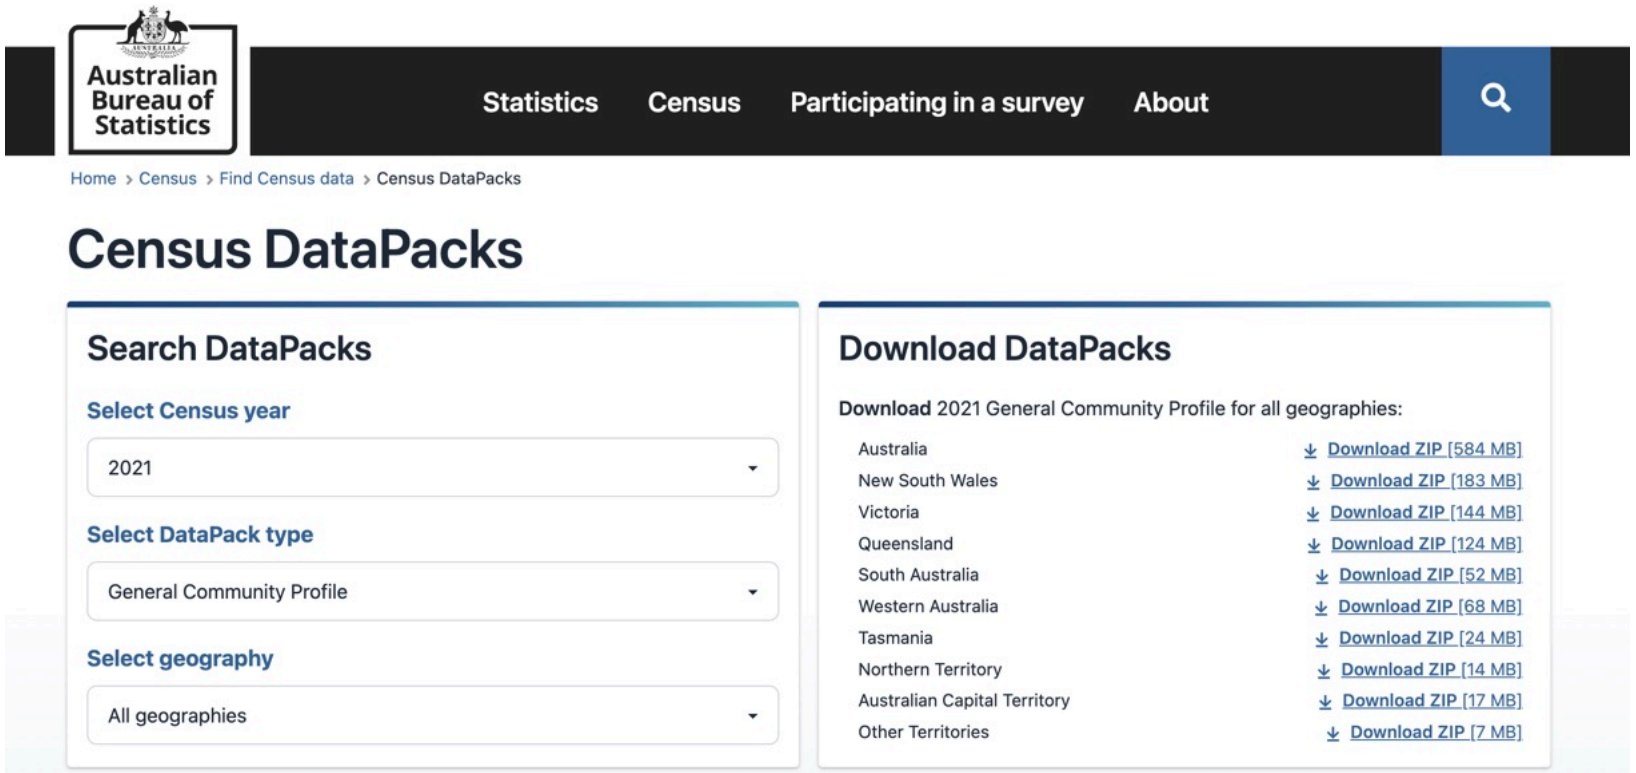

ABS Census 2021 GeoPackages

GeoPackage

> *A **GeoPackage** (GPKG) is an open, non-proprietary, platform-independent and standards-based data format for geographic information system implemented as a SQLite database container. Defined by the **Open Geospatial Consortium** (OGC) with the backing of the US military and published in 2014, GeoPackage has seen widespread support from various government, commercial, and open source organizations.*
>
> — *Wikipedia*

Recall: OGC also defines the WKT

ABS GeoPackage (2021)

https://datapacks.censusdata.abs.gov.au/geopackages/

1. 2021
2. Victoria
3. Topic: Employment and Income or Table: G17
4. GDA2020

- **Or use the `strayr` package!** We'll use the one from the ABS website instead.

```
# 2022
geopath_2021_G02 <- here::here("data/Geopackage_2021_G02_VIC_GDA2020/G02_VIC_GDA2020.gpkg")
st_layers(geopath_2021_G02)

## Driver: GPKG
## Available layers:
##            layer_name geometry_type features fields crs_name
## 1    G02_UCL_2021_VIC Multi Polygon      347     11  GDA2020
## 2    G02_SUA_2021_VIC Multi Polygon       22     11  GDA2020
## 3    G02_STE_2021_VIC Multi Polygon        1     11  GDA2020
## 4   G02_SOSR_2021_VIC Multi Polygon       12     11  GDA2020
## 5    G02_SOS_2021_VIC Multi Polygon        6     11  GDA2020
```

ABS GeoPackage (2016)

https://datapacks.censusdata.abs.gov.au/geopackages/

1. Victoria
2. Employment, Income and Unpaid Work (EIUW)
3. EIUW GeoPackage A

- **Or use the `strayr` package!** We'll use the one from the ABS website instead.

```
geopath_2016_eiuwa <- here::here("data/Geopackage_2016_EIUWA_for_VIC/census2016_eiuwa_vic_short.gpkg")
st_layers(geopath_2016_eiuwa)

## Driver: GPKG
## Available layers:
##                          layer_name geometry_type features fields crs_name
## 1    census2016_eiuwa_vic_ced_short                    39    489    GDA94
## 2  census2016_eiuwa_vic_gccsa_short                     4    489    GDA94
## 3    census2016_eiuwa_vic_lga_short                    82    489    GDA94
## 4    census2016_eiuwa_vic_poa_short                   698    489    GDA94
## 5     census2016_eiuwa_vic_ra_short                     6    489    GDA94
## 6    census2016_eiuwa_vic_sa1_short                 14073    489    GDA94
## 7    census2016_eiuwa_vic_sa2_short                   464    489    GDA94
```

The Australian Statistical Geography Standard (ASGS)

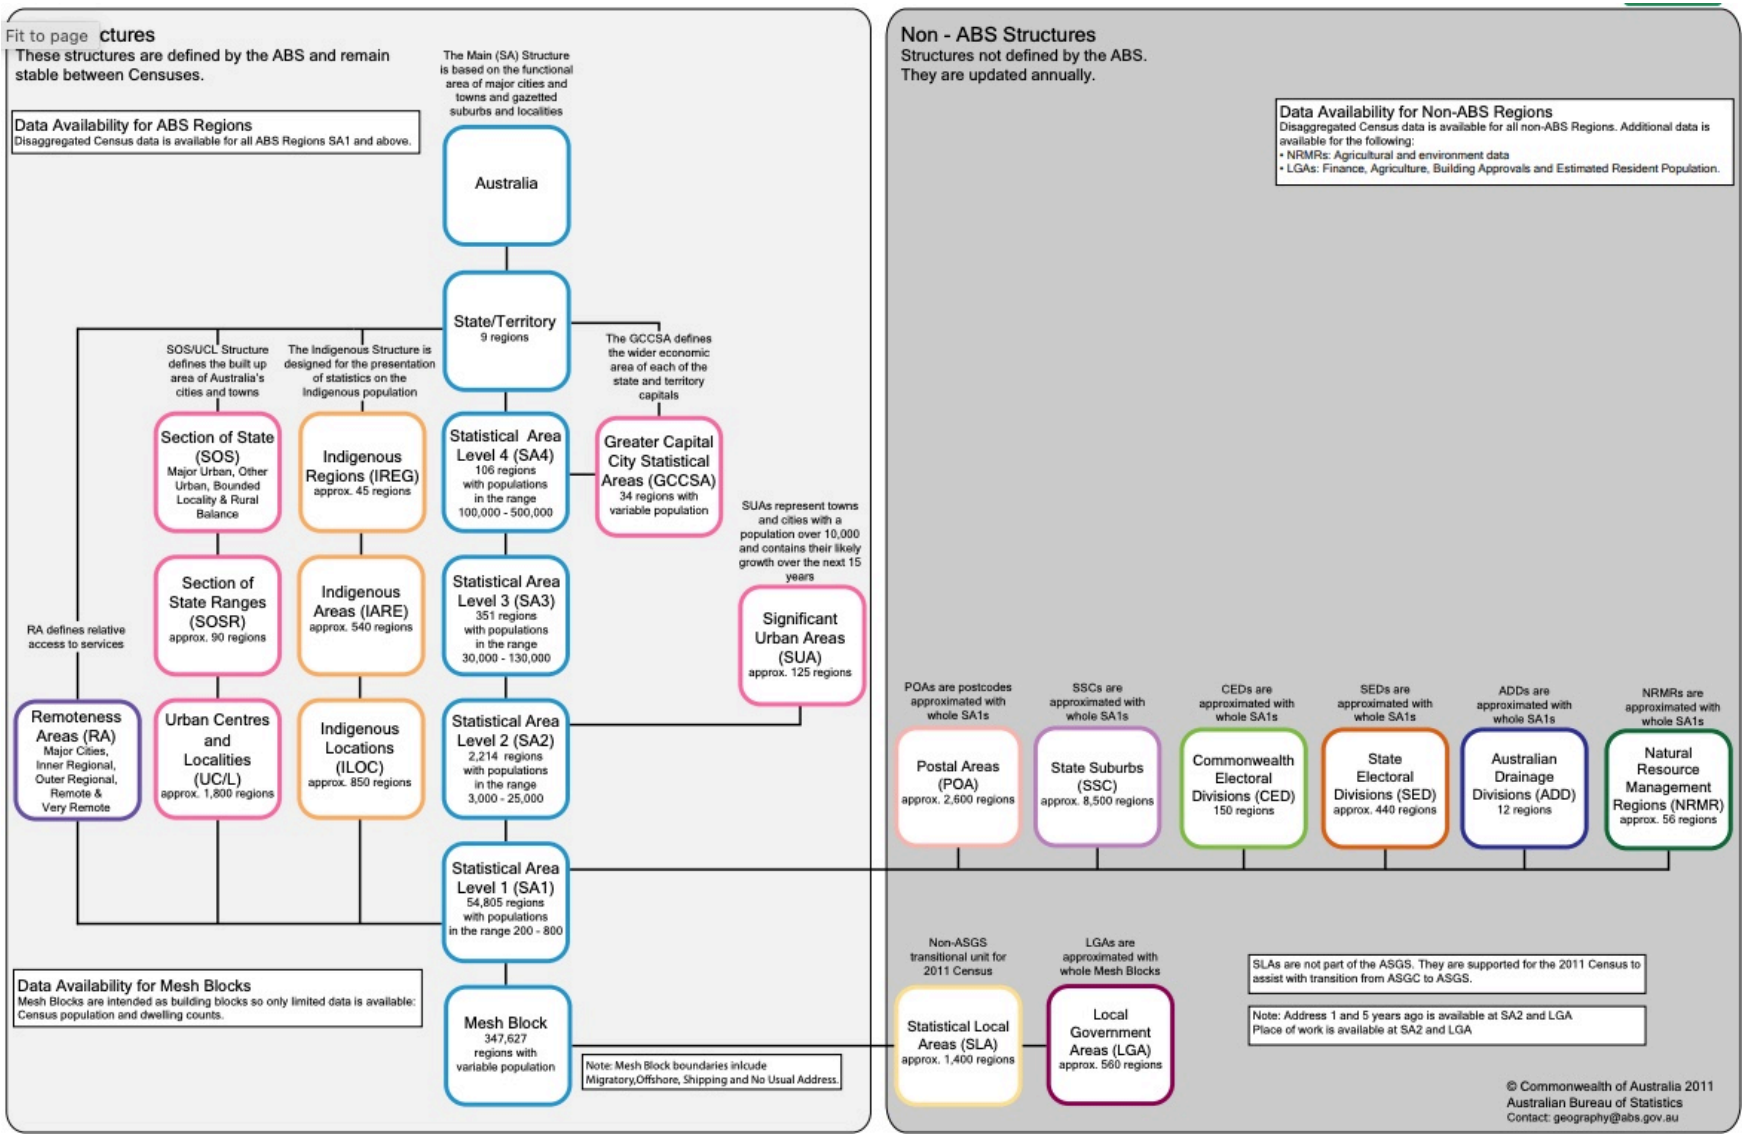

The number of regions for each layer

```
st_layers(geopath_2021_G02) %>%
  # make it into a data.frame first
  tibble(!!!.) %>%
  # then you can the dplyr operations
  dplyr::arrange(features)
```

```
## # A tibble: 16 × 6
##    name               geomtype  driver features fields crs
##    <chr>              <list>    <chr>     <dbl>  <dbl> <list>
##  1 G02_STE_2021_VIC   <chr [1]> GPKG          1     11 <crs>
##  2 G02_GCCSA_2021_VIC <chr [1]> GPKG          4     11 <crs>
##  3 G02_SOS_2021_VIC   <chr [1]> GPKG          6     11 <crs>
##  4 G02_RA_2021_VIC    <chr [1]> GPKG          6     11 <crs>
##  5 G02_SOSR_2021_VIC  <chr [1]> GPKG         12     11 <crs>
##  6 G02_SA4_2021_VIC   <chr [1]> GPKG         19     11 <crs>
##  7 G02_SUA_2021_VIC   <chr [1]> GPKG         22     11 <crs>
##  8 G02_CED_2021_VIC   <chr [1]> GPKG         41     11 <crs>
##  9 G02_SA3_2021_VIC   <chr [1]> GPKG         68     11 <crs>
## 10 G02_LGA_2021_VIC   <chr [1]> GPKG         82     11 <crs>
## 11 G02_SED_2021_VIC   <chr [1]> GPKG         90     11 <crs>
## 12 G02_UCL_2021_VIC   <chr [1]> GPKG        347     11 <crs>
## 13 G02_SA2_2021_VIC   <chr [1]> GPKG        524     11 <crs>
## 14 G02_POA_2021_VIC   <chr [1]> GPKG        694     11 <crs>
## 15 G02_SAL_2021_VIC   <chr [1]> GPKG       2946     11 <crs>
## 16 G02_SA1_2021_VIC   <chr [1]> GPKG      15482     11 <crs>
```

Data in the layer

```
vicmap_ste_G02 <- read_sf(geopath_2021_G02, layer = "G02_STE_2021_VIC")
vicmap_ste_G02$geom
```

```
## Geometry set for 1 feature
## Geometry type: MULTIPOLYGON
## Dimension:     XY
## Bounding box:  xmin: 140.9619 ymin: -39.15918 xmax: 149.9762 ymax: -33.98064
## Geodetic CRS:  GDA2020
```

```
str(vicmap_ste_G02)
```

```
## sf [1 × 12] (S3: sf/tbl_df/tbl/data.frame)
##  $ STE_CODE_2021               : chr "2"
##  $ STE_NAME_2021               : chr "Victoria"
##  $ Median_age_persons          : num 38
##  $ Median_mortgage_repay_monthly: num 1859
##  $ Median_tot_prsnl_inc_weekly  : num 803
##  $ Median_rent_weekly          : num 370
##  $ Median_tot_fam_inc_weekly    : num 2136
##  $ Average_num_psns_per_bedroom : num 0.8
##  $ Median_tot_hhd_inc_weekly    : num 1759
##  $ Average_household_size       : num 2.5
##  $ AREA_ALBERS_SQKM             : num 227496
##  $ geom                         :sfc_MULTIPOLYGON of length 1; first list element: List of 157
##   ..$ :List of 1
##   .. ..$ : num [1:39, 1:2] 146 146 146 146 146 ...
```

State or Territory (STE)

```
vicmap_ste_G02 <- read_sf(geopath_2021_G02, layer = "G02_STE_2021_VIC")

ggplot(vicmap_ste_G02) +
  geom_sf(aes(geometry = geom, fill = Median_age_persons))
```

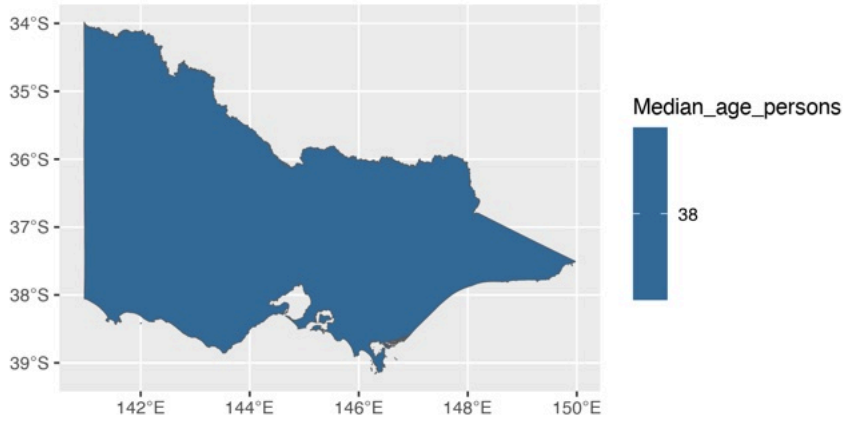


```
nrow(vicmap_ste)

## [1] 1
```

Breakout Session

Try it yourself time:

- Download the geopackage for the 2021 census and variable G02
- Look at how that spatial data is organised
- Look at the different layers

What are the differences between the regionalisations?

What do you notice?

Want a challenge visualise the 2016 data!

Greater Capital City Statistical Areas (GCCSA)

- Each region with variable population

```
vicmap_gccsa_G02 <- read_sf(geopath_2021_G02, layer = "G02_GCCSA_2021_VIC")
ggplot(vicmap_gccsa_G02) +
  geom_sf(aes(geometry = geom, fill = Median_age_persons))
```

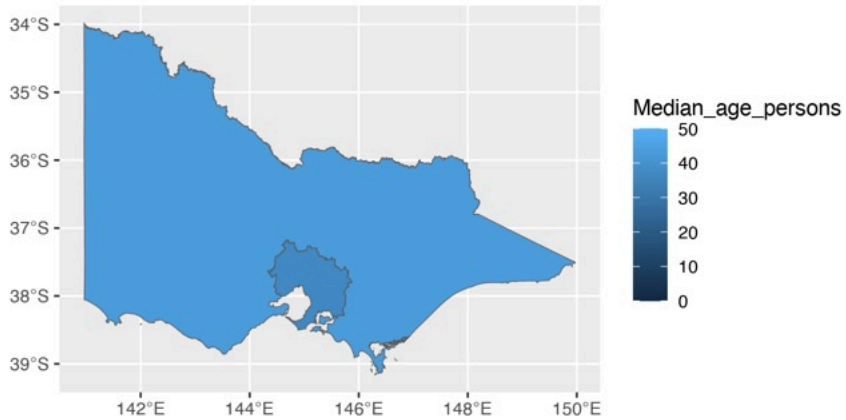


```
nrow(vicmap_gccsa_G02)
```

```
## [1] 4
```

Section of State (SOS)

- Major urban, other urban, bounded locally & rural balance

```
vicmap_sos_G02 <- read_sf(geopath_2021_G02, layer = "G02_SOS_2021_VIC")
ggplot(vicmap_sos_G02) +
  geom_sf(aes(geometry = geom, fill = Median_age_persons))
```

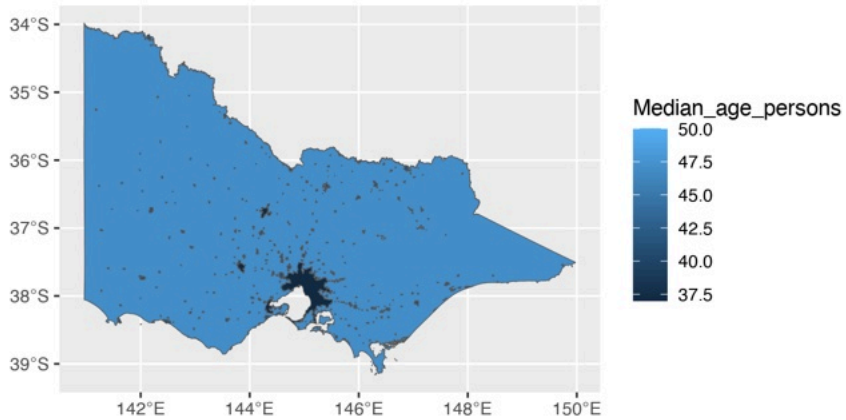


```
nrow(vicmap_sos_G02)
```

```
## [1] 6
```

Remoteness Areas (RA)

```
vicmap_ra_G02 <- read_sf(geopath_2021_G02, layer = "G02_RA_2021_VIC")
ggplot(vicmap_ra_G02) +
  geom_sf(aes(geometry = geom, fill = Median_age_persons))
```

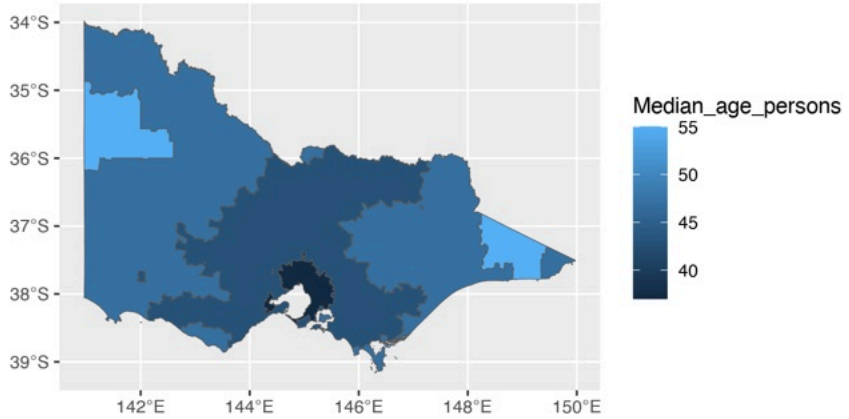


```
nrow(vicmap_ra_G02)

## [1] 6
```

Section of State Ranges (SOSR)

```
vicmap_sosr_G02 <- read_sf(geopath_2021_G02, layer = "G02_SOSR_2021_VIC")
ggplot(vicmap_sosr_G02) +
  geom_sf(aes(geometry = geom, fill = Median_age_persons))
```

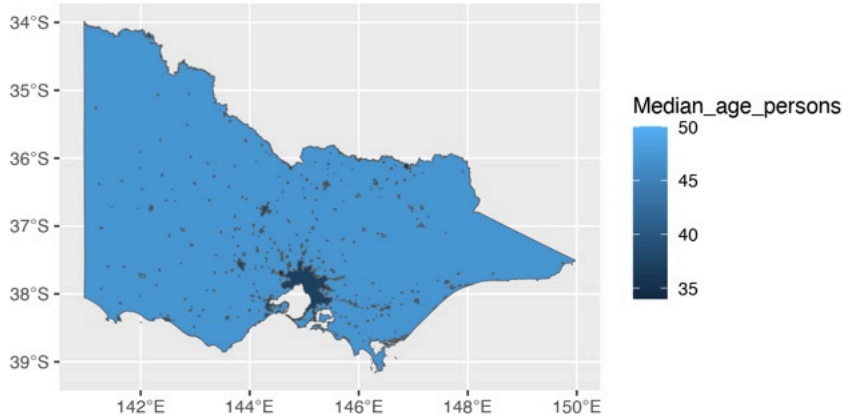


```
nrow(vicmap_sosr_G02)

## [1] 12
```

Statistical Area Level 4 (SA4)

- Each region with population of 100,000 - 500,000

```
vicmap_sa4_G02 <- read_sf(geopath_2021_G02, layer = "G02_SA4_2021_VIC")
ggplot(vicmap_sa4_G02) +
  geom_sf(aes(geometry = geom, fill = Median_age_persons))
```

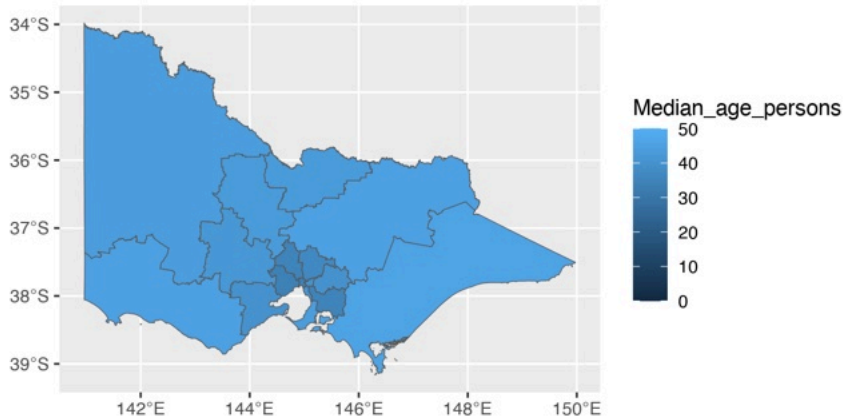


```
nrow(vicmap_sa4_G02)
```

```
## [1] 19
```

Significant Urban Areas (SUA)

```
vicmap_sua_G02 <- read_sf(geopath_2021_G02, layer = "G02_SUA_2021_VIC")
ggplot(vicmap_sua_G02) +
  geom_sf(aes(geometry = geom, fill = Median_age_persons))
```

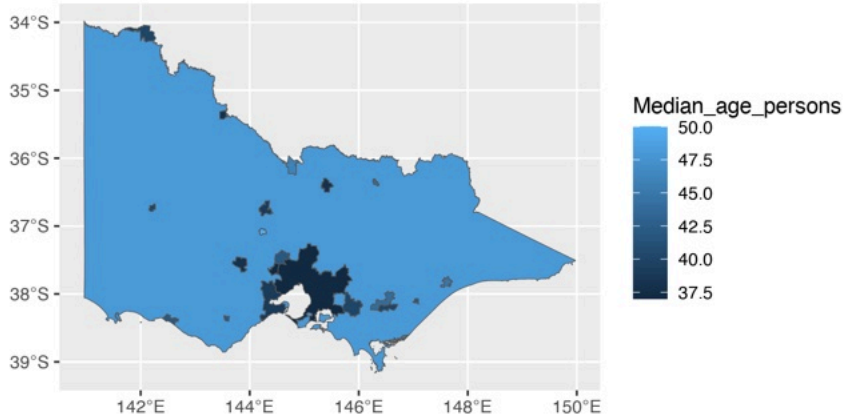


```
nrow(vicmap_sua_G02)

## [1] 22
```

Commonwealth Electoral Division (CED)

```
vicmap_ced_G02 <- read_sf(geopath_2021_G02, layer = "G02_CED_2021_VIC")
ggplot(vicmap_ced_G02) +
  geom_sf(aes(geometry = geom, fill = Median_age_persons))
```

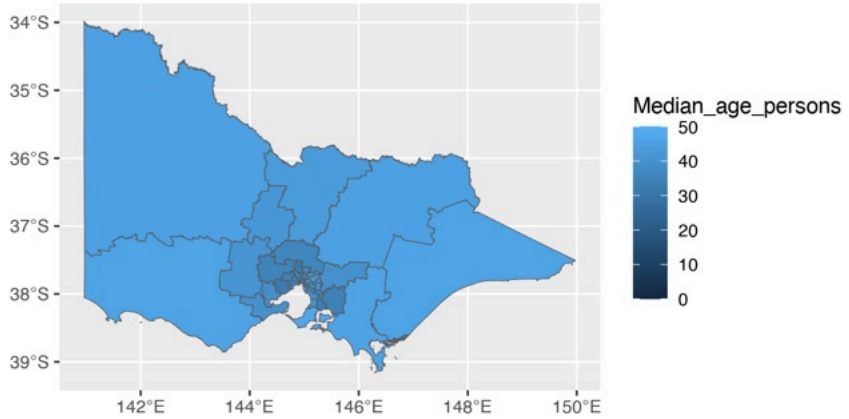


```
nrow(vicmap_ced_G02)

## [1] 41
```

Statistical Area Level 3 (SA3)

- Each region with population of 30,000 - 130,000

```
vicmap_sa3_G02 <- read_sf(geopath_2021_G02, layer = "G02_SA3_2021_VIC")
ggplot(vicmap_sa3_G02) +
  geom_sf(aes(geometry = geom, fill = Median_age_persons))
```

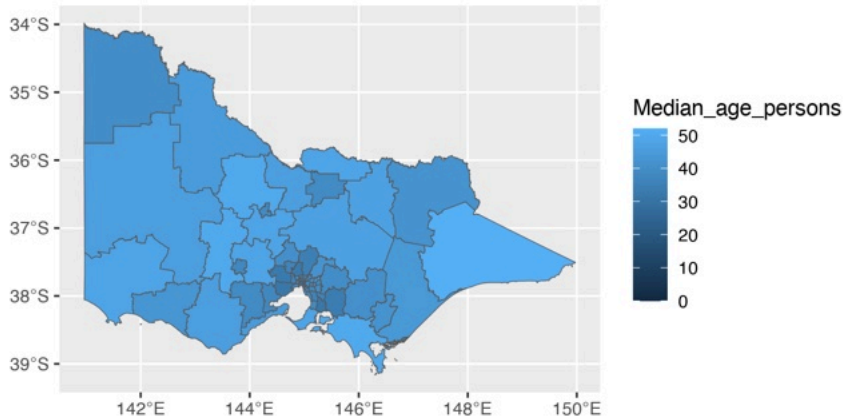


```
nrow(vicmap_sa3_G02)
```

```
## [1] 68
```

Local Government Area (LGA)

```
vicmap_lga_G02 <- read_sf(geopath_2021_G02, layer = "G02_SA3_2021_VIC")
ggplot(vicmap_lga_G02) +
  geom_sf(aes(geometry = geom, fill = Median_age_persons))
```

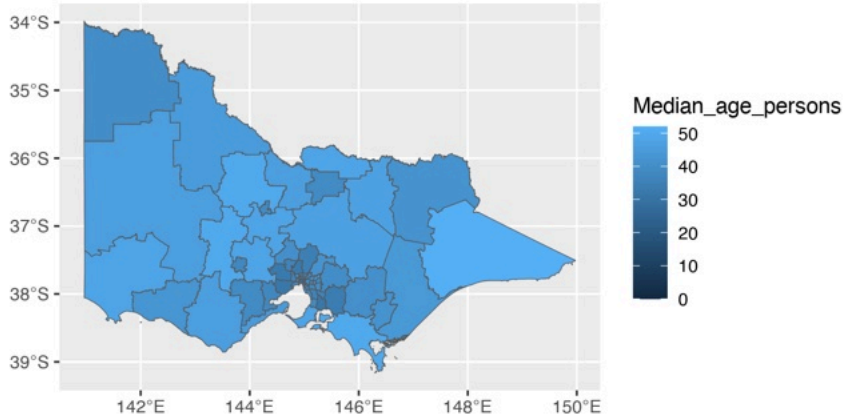


```
nrow(vicmap_lga_G02)

## [1] 68
```

State Electoral Division (SED)

```
vicmap_sed_G02 <- read_sf(geopath_2021_G02, layer = "G02_SED_2021_VIC")
ggplot(vicmap_sed_G02) +
  geom_sf(aes(geometry = geom, fill = Median_age_persons))
```

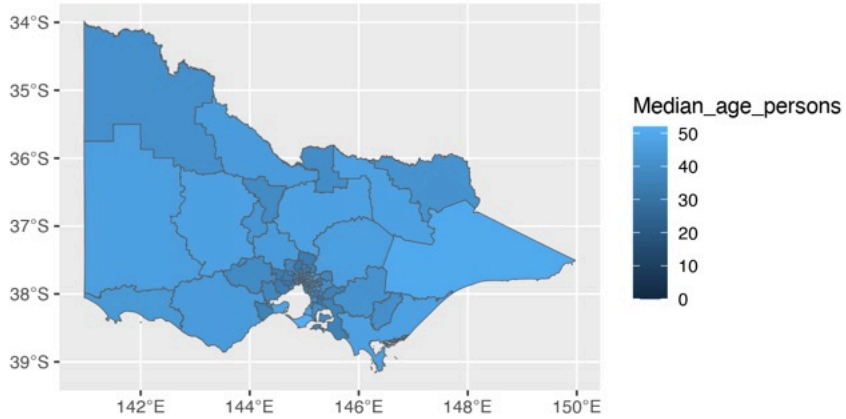


```
nrow(vicmap_sed_G02)

## [1] 90
```

Urban Centres and Localities (UCL)

```
vicmap_ucl_G02 <- read_sf(geopath_2021_G02, layer = "G02_UCL_2021_VIC")
ggplot(vicmap_ucl_G02) +
  geom_sf(aes(geometry = geom, fill = Median_age_persons))
```

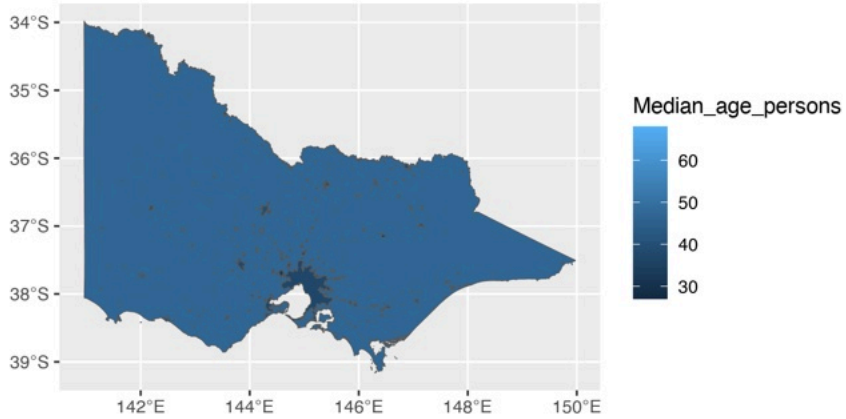


```
nrow(vicmap_ucl_G02)
```

```
## [1] 347
```

Statistical Area Level 2 (SA2)

- Each region with populations in the range of 3,000-25,000

```
vicmap_sa2_G02 <- read_sf(geopath_2021_G02, layer = "G02_SA3_2021_VIC")
ggplot(vicmap_sa2_G02) +
  geom_sf(aes(geometry = geom, fill = Median_age_persons))
```

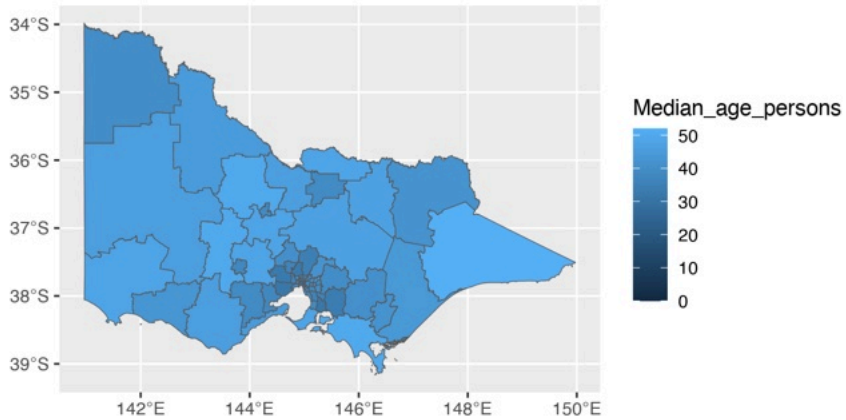


```
nrow(vicmap_sa2_G02)
```

```
## [1] 68
```

Postal Areas (POA)

```
vicmap_poa_G02 <- read_sf(geopath_2021_G02, layer = "G02_POA_2021_VIC")
ggplot(vicmap_poa_G02) +
  geom_sf(aes(geometry = geom, fill = Median_age_persons))
```

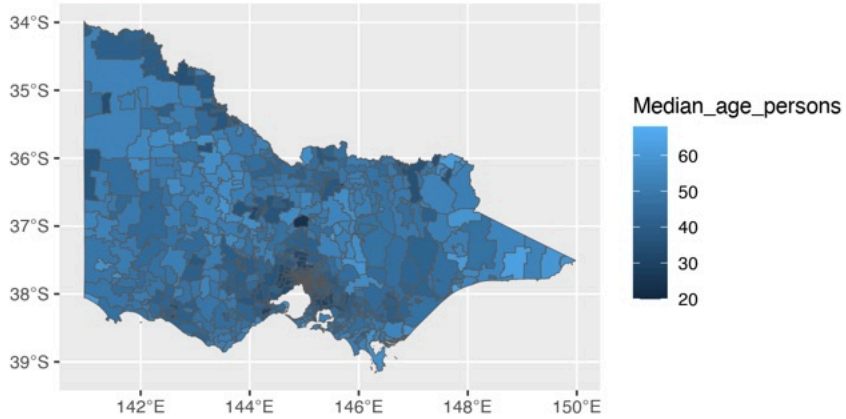


```
nrow(vicmap_poa_G02)

## [1] 694
```

State Area Localitites (SAL) (Formerly SSC)

```
vicmap_sal_G02 <- read_sf(geopath_2021_G02, layer = "G02_SAL_2021_VIC")
ggplot(vicmap_sal_G02) +
  geom_sf(aes(geometry = geom, fill = Median_age_persons))
```

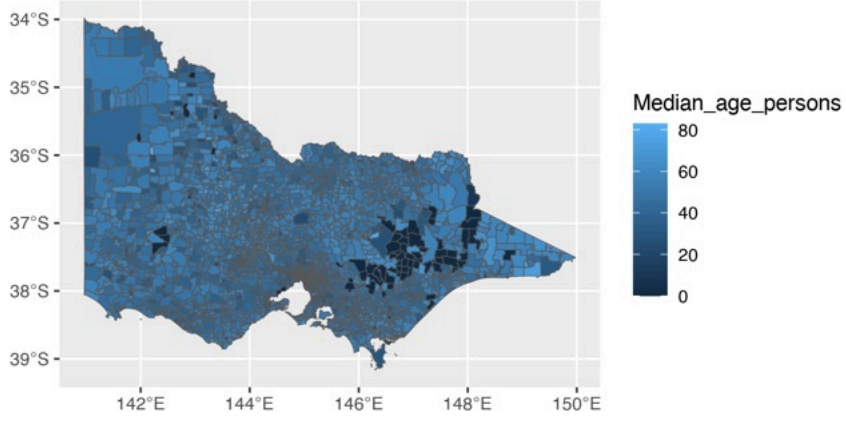


```
nrow(vicmap_sal_G02)

## [1] 2946
```

Statistical Area Level 1 (SA1)

- Each region with a population of range 200-800

```
vicmap_sa1_G02 <- read_sf(geopath_2021_G02, layer = "G02_SA1_2021_VIC")
ggplot(vicmap_sa1_G02) +
  geom_sf(aes(geometry = geom, fill = Median_age_persons))
```

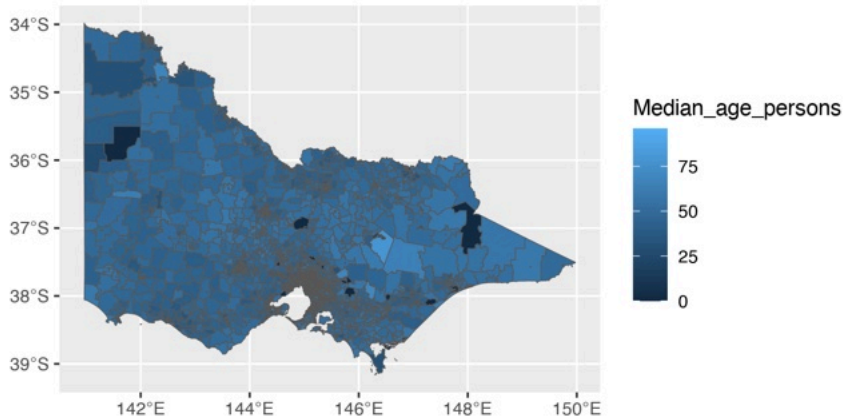


```
nrow(vicmap_sa1_G02)
```

```
## [1] 15482
```

Electorate boundary

vs

Census boundary

🎯 Estimate a median age for an electorate

Comparing SED 2021 and electorates divisions 2022

See here for `electoral_winners` data was.

```
ggplot() +
  geom_sf(data = vicmap_sed_G02,
          aes(geometry = geom, fill = Median_age_persons),
    alpha = 1, color = "white", size = 2) +
  geom_sf(data = winners_with_map, aes(geometry = geometry),
    fill = "transparent", color = "red", size = 2) +
  coord_sf(xlim = c(144.95, 145.24), ylim = c(-38.05, -37.85)) +
  theme_bw()
```

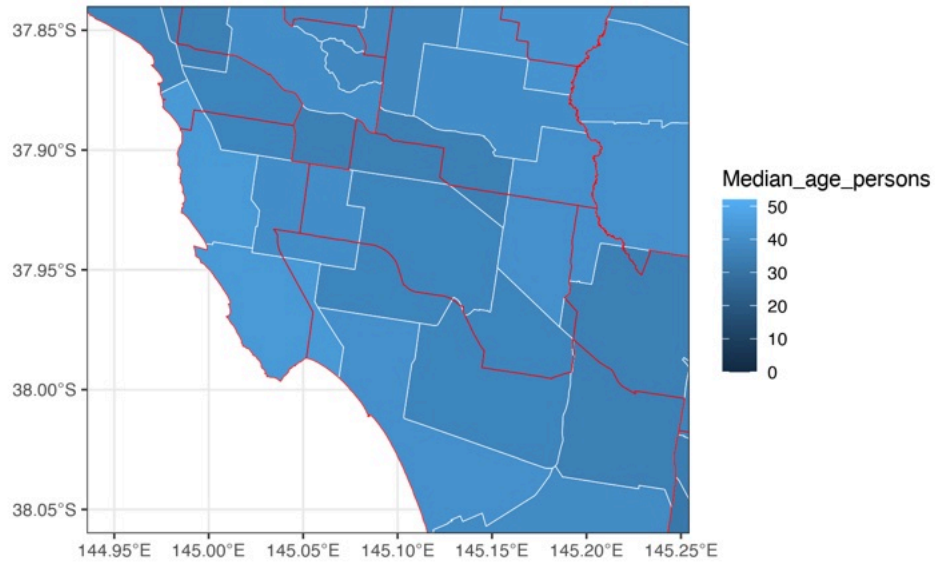

Closer look Hotham electorate ❶

```
electorate <- winners_with_map |>
  filter(DivisionNm == "HOTHAM")

# Set projection to GDA1994 using EPSG:4283
st_crs(electorate$geometry,4283)

# Transform projection from GDA1994 to GDA2020 using EPSG:7844
electorate$geometry = st_transform(electorate$geometry, 7844)

sed_intersect <- vicmap_sed_G02 |>
  filter(st_intersects(geom,
    electorate$geometry,
    sparse = FALSE
  )[, 1])

ggplot() +
  geom_sf(data = sed_intersect,
    aes(geometry = geom), color = "red", fill = "transparent") +
  geom_sf_text(data = sed_intersect,
    aes(label = SED_CODE_2021, geometry = geom), color = "red") +
  geom_sf(data = electorate, aes(geometry = geometry), fill = "transparent") +
  geom_sf_text(data = electorate, aes(geometry = geometry, label = DivisionNm))
```

```
## Coordinate Reference System:
##   User input: GDA94
##   wkt:
```

Closer look Hotham electorate 1

There are 7 SED regions that intersect with Hotham electorate.

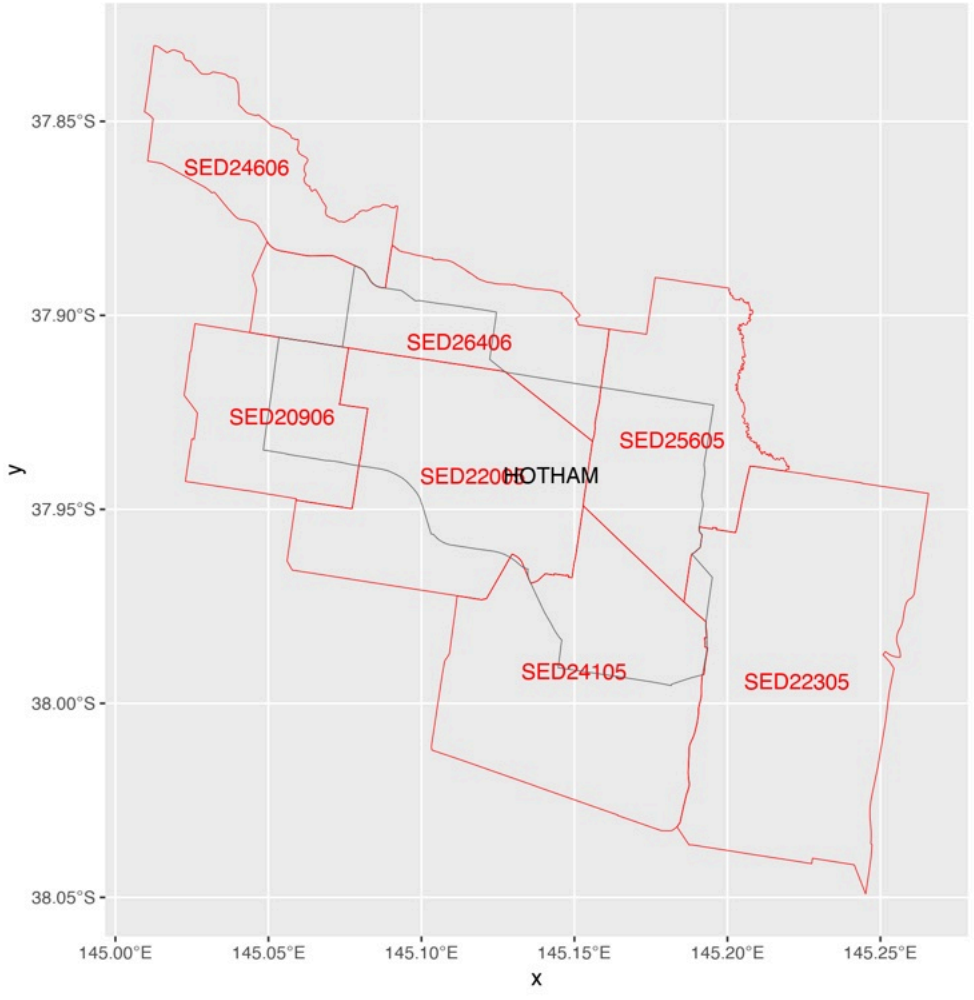

Closer look Hotham electorate 2

```
sed_intersect2 <- sed_intersect |>
  mutate(
    geometry = st_intersection(geom, electorate$geometry),
    perc_area = 100 * st_area(geometry) / st_area(geom),
    perc_area = as.numeric(perc_area)
  ) |>
  filter(perc_area > 5)

ggplot(sed_intersect2, aes(geometry = geometry)) +
  geom_sf(data = electorate) +
  geom_sf_text(
    data = electorate,
    aes(label = DivisionNm)
  ) +
  geom_sf(color = "red", aes(fill = Median_age_persons)) +
  geom_sf_text(
    aes(
      label = glue::glue("{SED_CODE_2021}
                         ({scales::comma(perc_area, 1)}%, {Median_age_persons})")
    ),
    color = "red"
  ) +
  theme(legend.position = "bottom")
```

Closer look Hotham electorate ❷

- There are 5 SED areas with at least 5% intersection with the electoral area.
- **How would you characterise the median age for Hotham?**

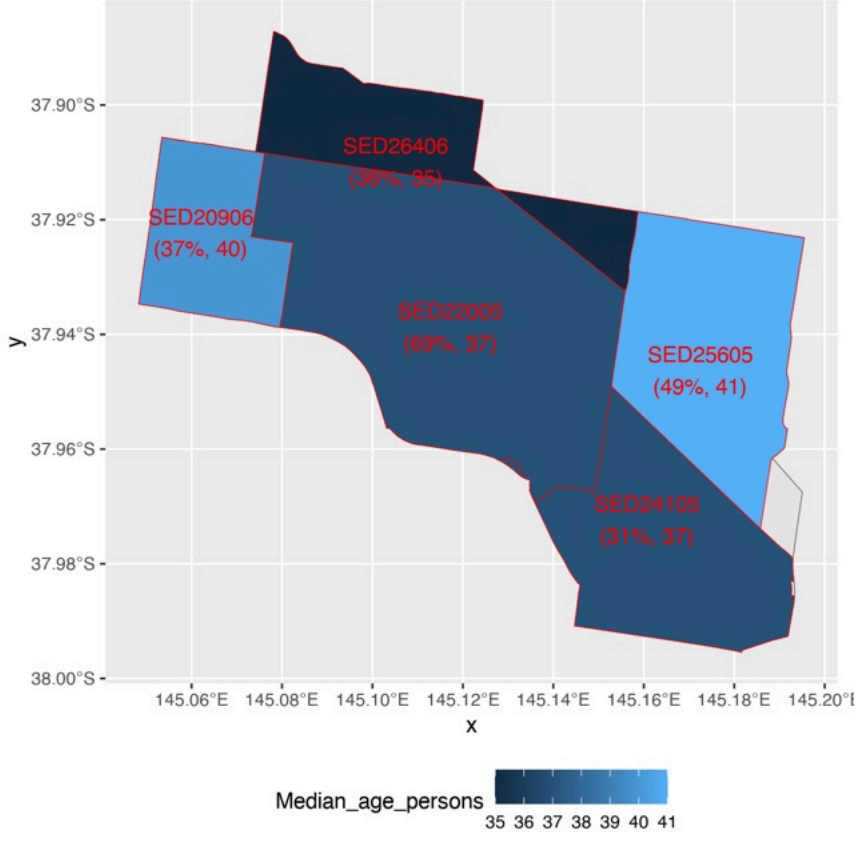

Closer look Hotham electorate 3

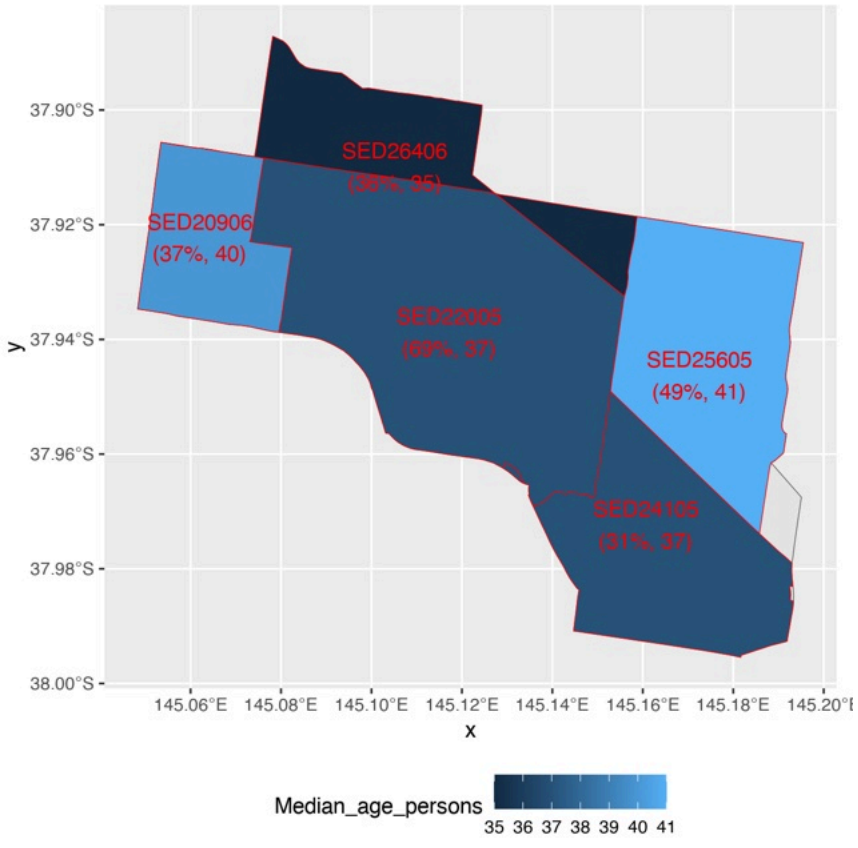

Strategy 1

```
sort(sed_intersect2$Median_age_persons)
```

```
## [1] 35 37 37 40 41
```

Strategy 2

```
mean(sed_intersect2$Median_age_persons)
```

```
## [1] 38
```

Strategy 3

```
weighted.mean(
  sed_intersect2$Median_age_persons,
  sed_intersect2$perc_area
)
```

```
## [1] 38.06118
```

Closer look Hotham electorate 4

```
sa1_intersect <- vicmap_sa1_G02 %>%
  filter(st_intersects(geom,
    electorate$geometry,
    sparse = FALSE
  )[, 1])

sa1_intersect2 <- sa1_intersect %>%
  mutate(
    geometry = st_intersection(geom, electorate$geometry),
    perc_area = 100 * st_area(geometry) / st_area(geom),
    perc_area = as.numeric(perc_area)
  ) %>%
  filter(perc_area > 5)

ggplot(sa1_intersect) +
  geom_sf(color = "red",
          aes( fill = Median_age_persons, geometry = geom)) +
  geom_sf(data = electorate,
          color = "white", size = 2, fill = "transparent",
          aes(geometry = geometry)) +
  theme(legend.position = "bottom")
```

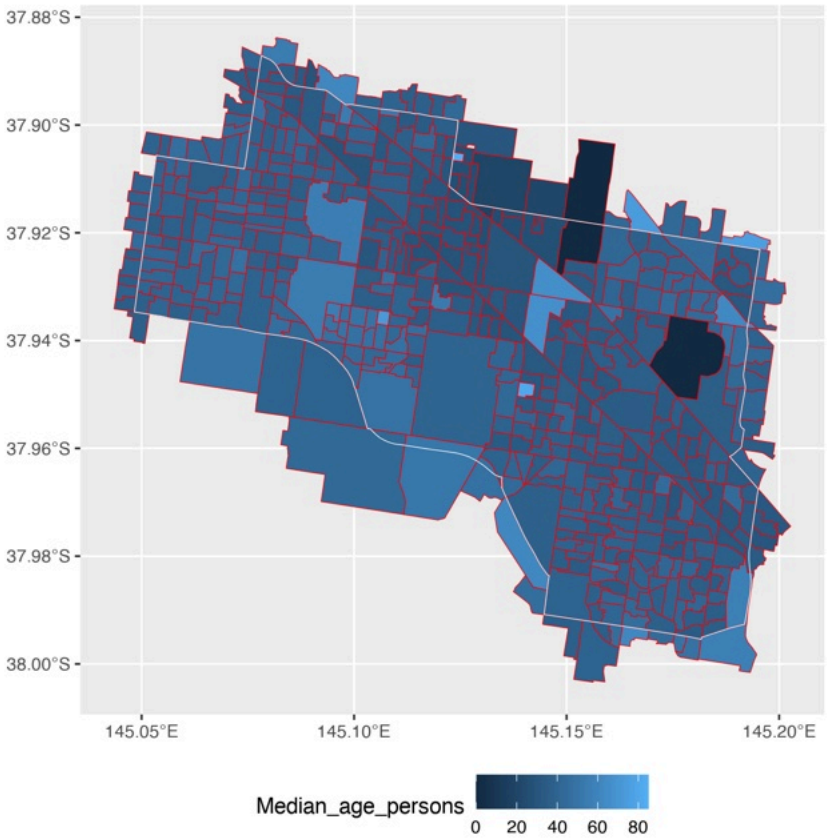

Closer look Hotham electorate 5

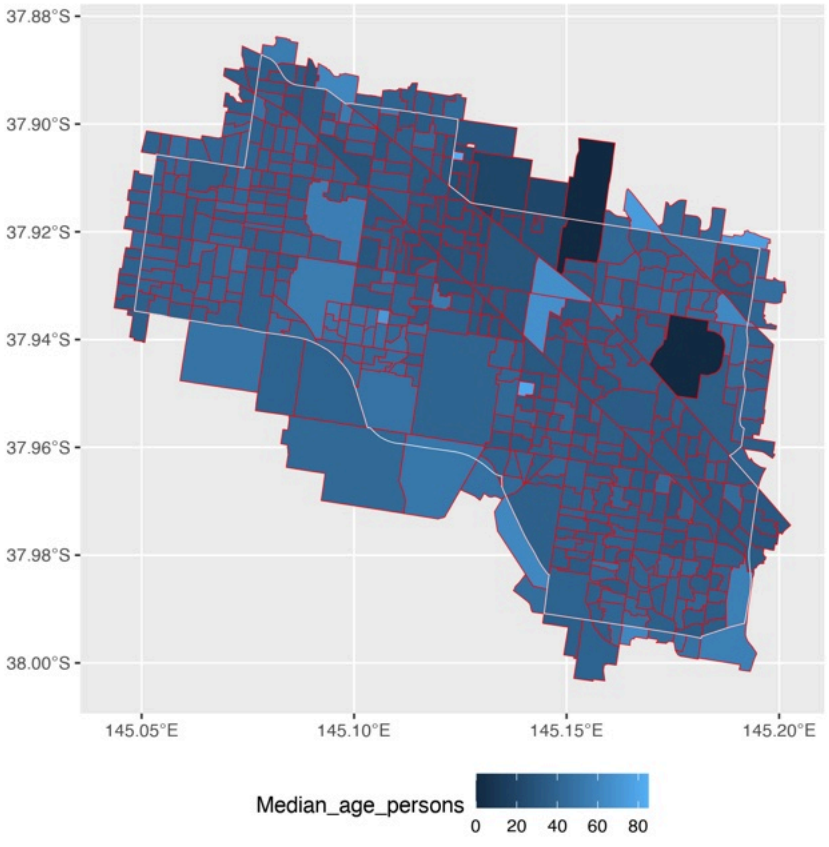

Strategy 1

```
fivenum(sa1_intersect2$Median_age_persons)

## [1]  0 35 38 41 81
```

Strategy 2

```
mean(sa1_intersect2$Median_age_persons)

## [1] 38.07674
```

Strategy 3

```
weighted.mean(sa1_intersect2$Median_age_persons, sa1_intersect2$perc_are

## [1] 38.02792
```

Strategy 4

```
ggplot(sa1_intersect2, aes(x = Median_age_persons)) +
  geom_histogram(binwidth = 1)
```

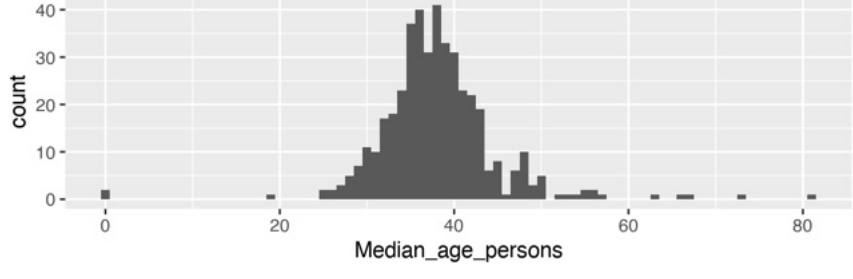

Closer look Zero median age

(Hotham 2022)[https://www.aec.gov.au/profiles/vic/hotham.htm]

```
sa1_intersect2 %>%
  filter(Median_age_persons == 0) %>%
  ggplot() +
  geom_sf() +
  geom_sf(
    data = electorate, color = "red",
    fill = "transparent",
    aes(geometry = geometry)
  )
```

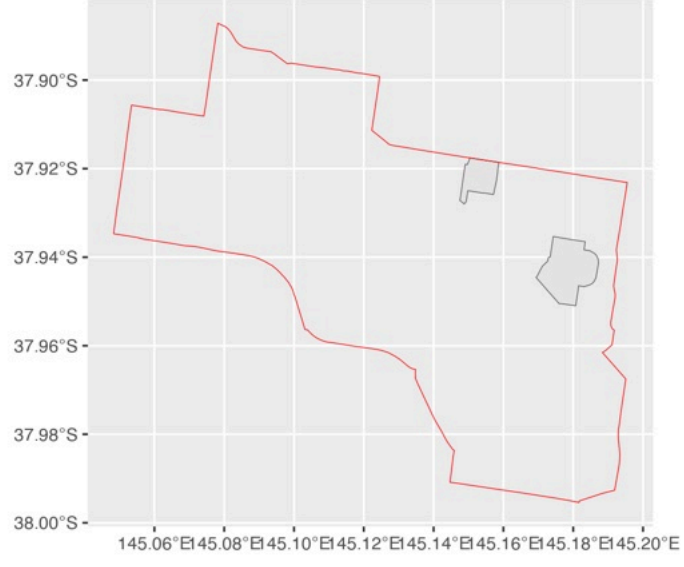

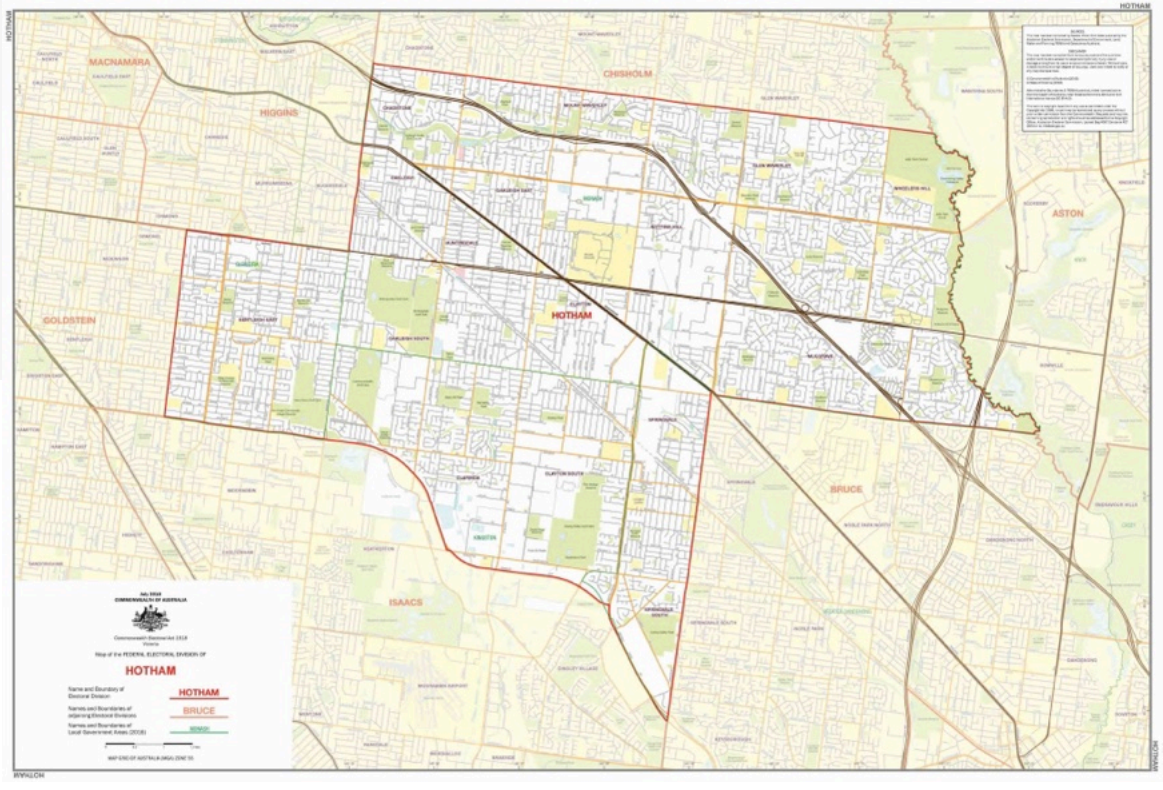

Closer look Hotham electorate 6

Before

Strategy 1

```
fivenum(sa1_intersect2$Median_age_persons)
```

```
## [1]  0 35 38 41 81
```

Strategy 2

```
mean(sa1_intersect2$Median_age_persons)
```

```
## [1] 38.07674
```

Strategy 3

```
weighted.mean(sa1_intersect2$Median_age_persons, sa1_interse
```

```
## [1] 38.02792
```

After

```
sa1_intersect3 <- sa1_intersect2 %>%
  filter(Median_age_persons != 0)
```

Strategy 1

```
fivenum(sa1_intersect3$Median_age_persons)
```

```
## [1] 19 35 38 41 81
```

Strategy 2

```
mean(sa1_intersect3$Median_age_persons)
```

```
## [1] 38.25467
```

Strategy 3

```
weighted.mean(sa1_intersect3$Median_age_persons, sa1_inter
```

```
## [1] 38.20661
```

Dorling Cartogram

```
sa1_intersect4 <- sa1_intersect %>%
  mutate(centroid = st_centroid(geom))

dorling_plot <- ggplot(sa1_intersect4) +
  geom_sf(
    data = electorate,
    aes(geometry = geometry), size = 4, fill = "grey60"
  ) +
  geom_sf(aes(geometry = centroid, color = Median_age_persons),
    size = 0.5, shape = 3
  ) +
  scale_color_viridis_c(name = "Median age", option = "magma")
```

Closer look Hotham electorate 7

```
sa1_intersect5 <- sa1_intersect4 %>%
  filter(st_intersects(centroid, electorate$geometry, sparse = FALSE)[, 1],
    Median_age_persons != 0)
```

Strategy 1

```
fivenum(sa1_intersect5$Median_age_persons)
```

```
## [1] 19 35 38 41 81
```

Strategy 2

```
mean(sa1_intersect5$Median_age_persons)
```

```
## [1] 38.20706
```

Strategy 4

```
ggplot(sa1_intersect5, aes(x = Median_age_persons)) +
  geom_histogram(binwidth = 1)
```

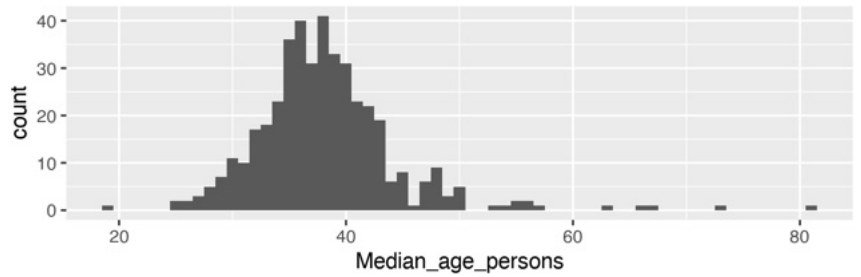

Summary

- There are many ways to characterise an electorate.
- Estimates of median age of an electorate is more consistent using SA1 map data than SED map data.

- We looked at mapping the 2021 census boundaries and projected a summary of the census variable (i.e. median age) onto a 2022 electoral district
- Discovered some of the challenges with matching two different types of data

Read Forbes, Cook & Hyndman (2020) Spatial modelling of the two-party preferred vote in Australian federal elections: 2001–2016. *Australian & New Zealand Journal of Statisitcs*. for a more sophisticated approach to studying the census variables and election results together.

Slides updated and maintained by Dr. Kate Saunders. Slides originally developed by Dr. Emi Tanaka

Lecturer: *Kate Saunders*

Department of Econometrics and Business Statistics

ETC5512.Clayton-x@monash.edu

Week 6